

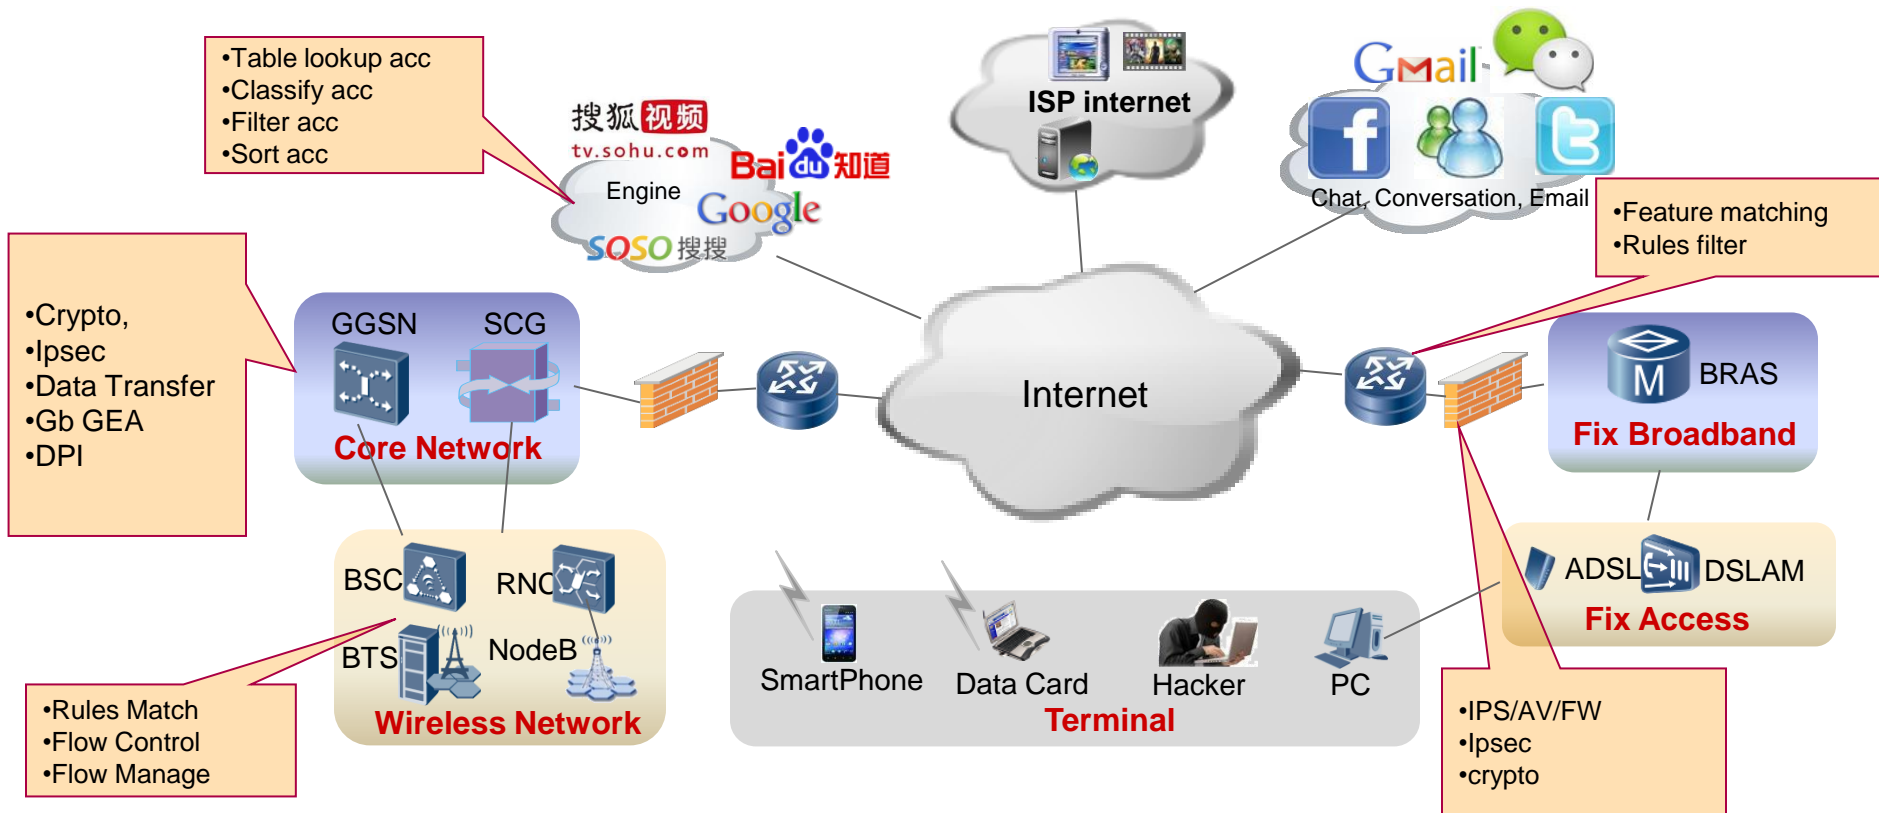
A new framework of cryptography virtio driver

Lei Gong

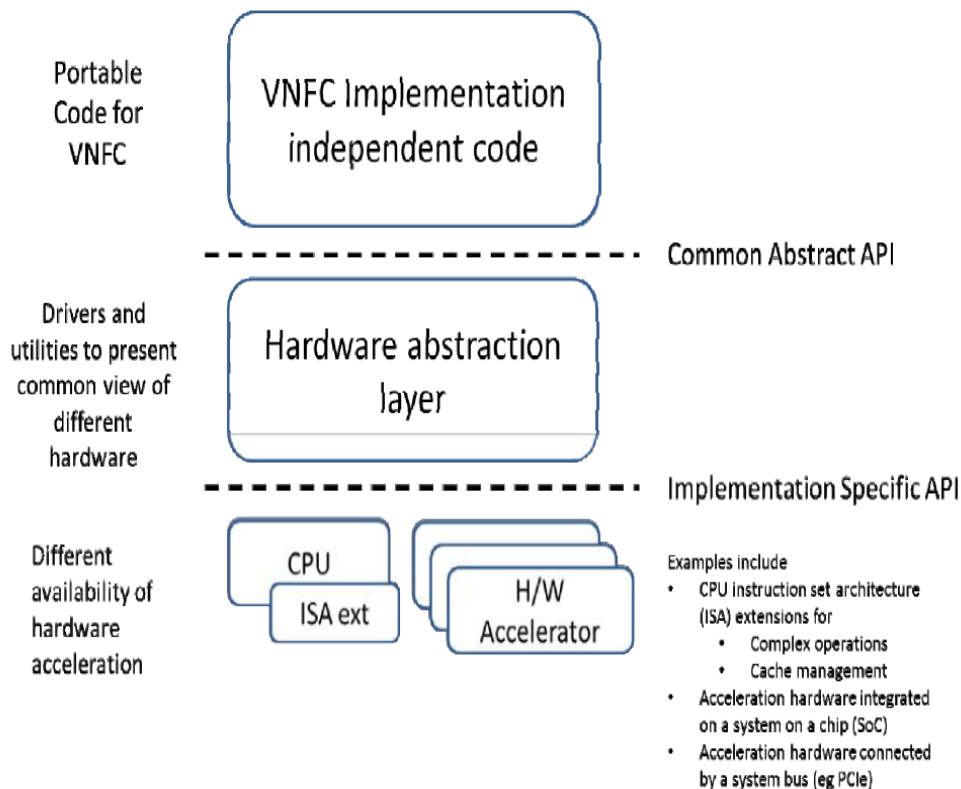
arei.gonglei@huawei.com

www.huawei.com

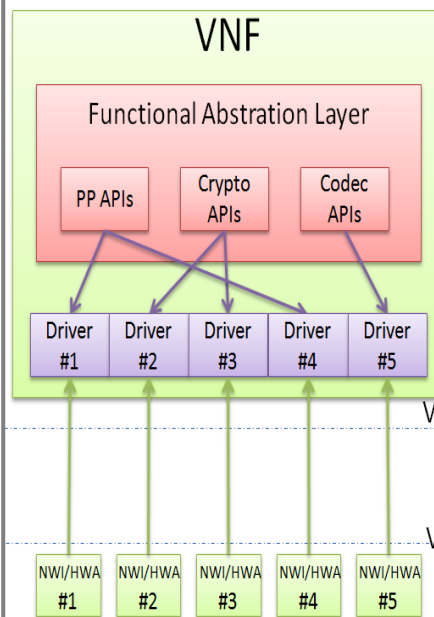
Scenarios of Hardware Acceleration



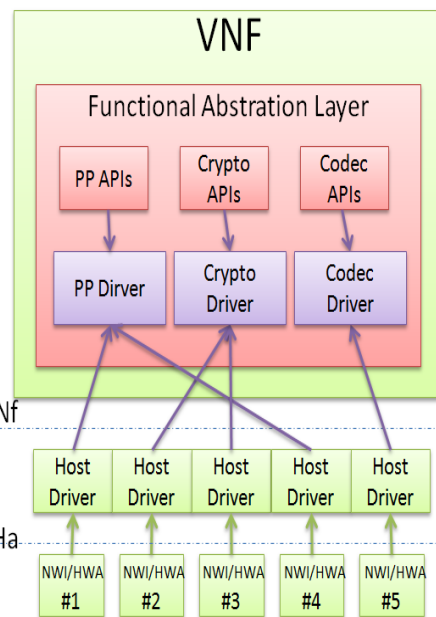
Abstract model of Accelerators (NFV)



"Pass-Through" Model



"Fully Intermediated" Model



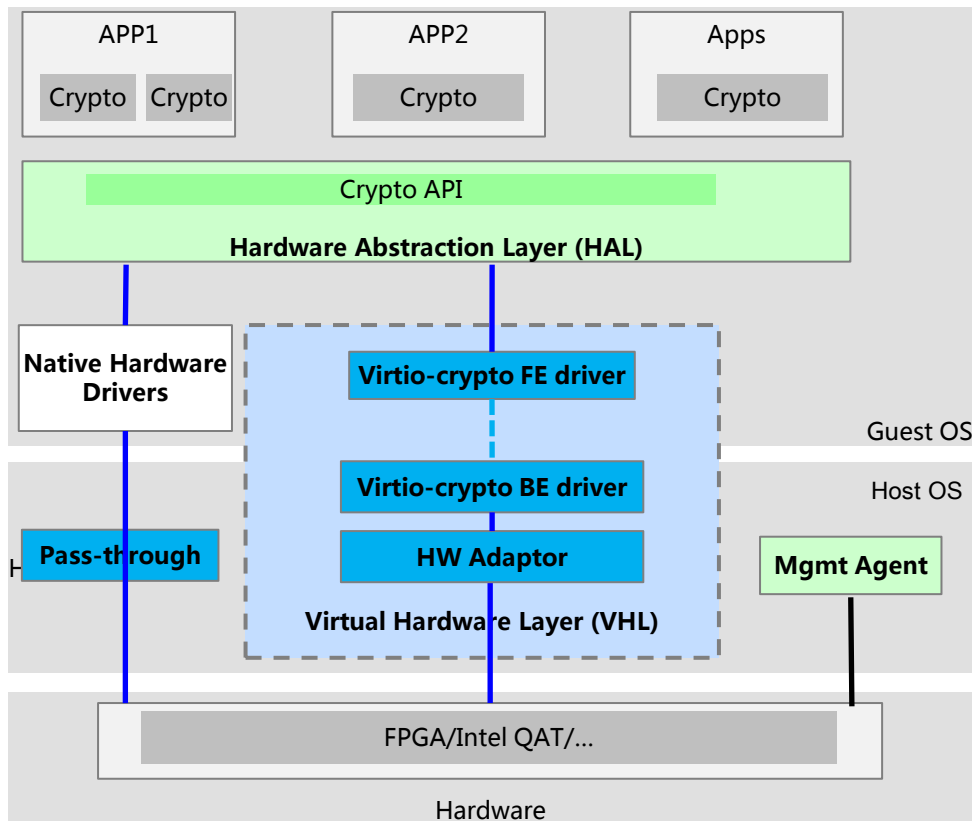
https://wiki.opnfv.org/dpacc/dpacc_project_proposal

Note: the original figure forwarded from ETSI GS NFV-INF 003 V1.1.1 (2014-12)

Why Virtio-crypto?

- Programmability
- Portability
- Scalability
- Hardware agnostic

Virtualization of Cryptography Accelerator



1. HAL

Provide acceleration APIs and runtimes

2. VHL

Provide virtual accelerators:

- 1) virtio-crypto FE driver
- 2) virtio-crypto BE driver
- 3) HW Adaptor : support different crypto accelerators

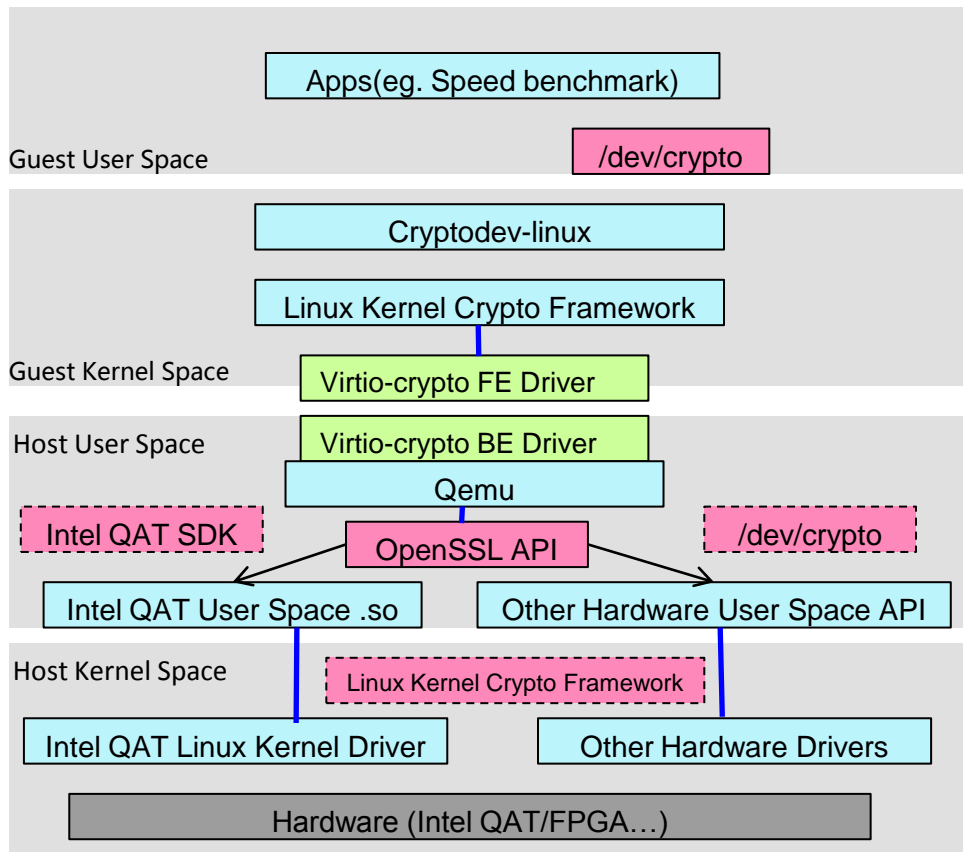
3. Pass-through

Accelerator pass-through

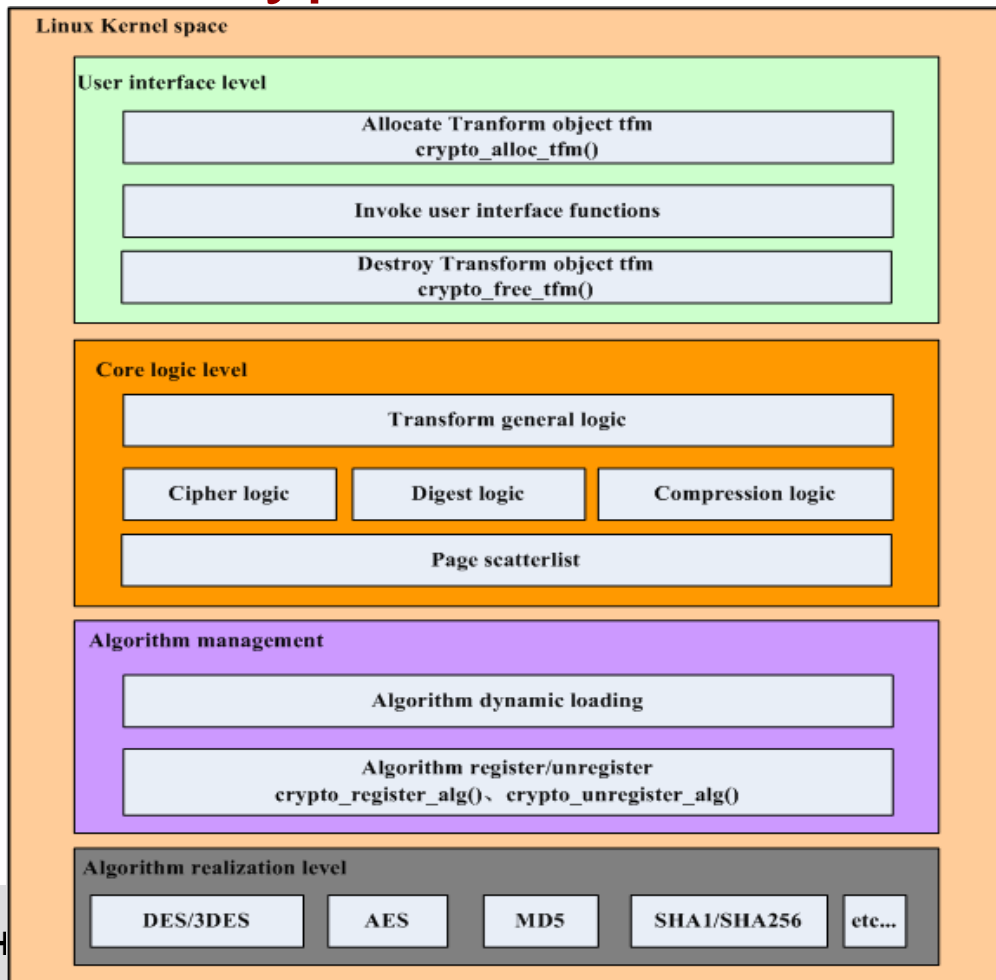
4. Mgmt Agent

Accelerator management

Flow of Virtio-crypto Prototype



Linux Kernel Crypto Framework



The kernel Crypto API

- ◆ A cryptography framework in the Linux kernel
- ◆ Can do Cipher, Hash, Compress, RNG, . . .
- ◆ Used by:
 - ✓ Network stack: IPsec, . . .
 - ✓ Device Mapper: dm-crypt, RAID, . . .
 - ✓ Userland Accessing:
 - ✓ AF_ALG
 - ✓ Cryptodev
- ◆ Maillist: linux-crypto@vger.kernel.org

AF_ALG introduction

- * Supports CIPHER, HASH
- * Socket-based interface
- + In-kernel code for years
- + Inherently asynchronous
- OpenSSL has out-of-tree engine for AF ALG
- GnuTLS does not have support for AF ALG
- Not many examples
- Higher latency

Cryptodev introduction

- * Supports CIPHER, HASH, AEAD
- * Uses character device interface
- + Compatible with OpenBSD /dev/crypto
- + API compatible, not OpenBSD code
- + OpenSSL has engine for cryptodev
- + GnuTLS has support for cryptodev
- + Has nice examples
- + Lower latency
- Out of kernel tree code (for years)
- Adds arbitrary IOCTLs

Cryptodev howto

Cryptodev usage pattern:

- a) `int cfd = open("/dev/crypto");`
- b) Fill in common struct cryptodev ctx
- c) Fill in struct crypt op
- d) Pass struct crypt op into kernel via `ioctl()`
- e) Retrieve results
- f) `close(cfd);`

Virtio-crypto BE driver

- * Emulate virtio-crypto devices in Qemu:
 Command line: `-device virtio-crypto-pci,id=crypto0`
- * Support different backend drivers:
 OpenSSL, Cryptodev, Intel QAT SDK
- * Support multiple virtio devices for each VM
- * Fit Virtio-1.0 spec
- * Cooperate with the virtio-crypto driver in guest

Virtio-crypto device

```
# lspci -v  
[skip]  
00:05.0 Unclassified device [00ff]: Red Hat, Inc Device 103f  
    Subsystem: Red Hat, Inc Device ffff  
    Flags: bus master, fast devsel, latency 0, IRQ 34  
    I/O ports at c000 [size=512]  
    Memory at febd3000 (32-bit, non-prefetchable) [size=4K]  
    Capabilities: [40] MSI-X: Enable+ Count=2 Masked-  
    Kernel driver in use: virtio-pci  
    Kernel modules: virtio_pci
```

Virtio-crypto FE driver

- * As a hardware crypto device
- * Support different algorithms:
Cipher, Hash, AEAD
- * Support multiple virtio devices for each VM
- * Fit Virtio-1.0 spec

Virtio-crypto module

```
# modinfo virtio-crypto
filename:      virtio-crypto.ko
author:        Gonglei <arei.gonglei@huawei.com>
license:       GPL
description:    Virtio crypto device driver
srcversion:    B5B95C74287DAE3AB7C134D
alias:         virtio:d0000FFFFv*
depends:        virtio_ring,virtio
vermagic:      3.0.76-0.11 SMP mod_unload modversions
parm:          virtio_crypto_verbosity:0: normal, 1: verbose, 2:
debug (int)
```

Register algorithms

```
static struct crypto_alg virtio_crypto_algs[] = { {  
    .cra_name = "cbc(aes)",  
    .cra_driver_name = "virtio_crypto_aes_cbc",  
    .cra_priority = 4001,  
    .cra_flags = CRYPTO_ALG_TYPE_ABLKCIPHER | CRYPTO_ALG_ASYNC,  
    .cra_blocksize = AES_BLOCK_SIZE,  
    .cra_ctxsize = sizeof(struct virtio_crypto_ablkcipher_ctx),  
    .cra_alignmask = 0,  
    .cra_module = THIS_MODULE,  
    .cra_type = &crypto_ablkcipher_type,  
    .cra_init = virtio_crypto_ablkcipher_init,  
    .cra_exit = virtio_crypto_ablkcipher_exit,  
    .cra_u = {  
        .ablkcipher = {  
            .setkey = virtio_crypto_ablkcipher_setkey,  
            .decrypt = virtio_crypto_ablkcipher_decrypt,  
            .encrypt = virtio_crypto_ablkcipher_encrypt,  
            .min_keysize = AES_MIN_KEY_SIZE,  
            .max_keysize = AES_MAX_KEY_SIZE,  
            .ivsize = AES_BLOCK_SIZE,  
        },  
    },  
},...
```


Virtio-crypto synchronize running


```
#cryptodev-linux-1.7 # ./tests/cipher -  
requested cipher CRYPTO_AES_CBC, got cbc(aes) with driver virtio_crypto_aes_cbc  
AES Test passed  
requested cipher CRYPTO_AES_CBC, got cbc(aes) with driver virtio_crypto_aes_cbc  
requested cipher CRYPTO_AES_CBC, got cbc(aes) with driver virtio_crypto_aes_cbc  
Test passed
```

Virtio-crypto asynchronize running

```
#cryptodev-linux-1.7 # ./tests/async_cipher -  
cryp1 written out  
cryp2 written out  
cryp1 + cryp2 successfully read  
result 1 passed  
result 2 passed  
AES Test passed  
running test_crypto  
test_crypto: got the session  
test_crypto: data encrypted  
test_crypto: session finished  
test_crypto: got new session  
test_crypto: data encrypted  
Test passed
```

Performance

Crypto-dev speed/async-speed benchmark (MB/sec) AES-128-CBC



- **Hardware**

- 1) Intel(R) Xeon(R) CPU E5-2620 v3 @ 2.40GHz

- 2) Intel QAT Coletto Creek PCIe DH895xCC SKU2

- **Software**

- Guest: Suse11.3 with 8 GB memory, 8vcpu

- Host: KVM 3.12, QEMU 2.4-rc3

TODO:

1. Performance optimization:
virtio-crypto-dataplane, batch processing, etc.
2. Other crypto algorithms support
3. Virtio-crypto upstream:
virtio-crypto spec, virtio-crypto code...

Q/A

Thank you!