Introduction and Improvement of PSI

Chengming Zhou
ByteDance STE Team

ByteDance 字节跳动

CLK 2021

Agenda

- Background
- PSI introduction
- PSI improvement
- Status & Future Work

Background

ByteDance 字节跳动

Resource

Maximizing Resource Utilization

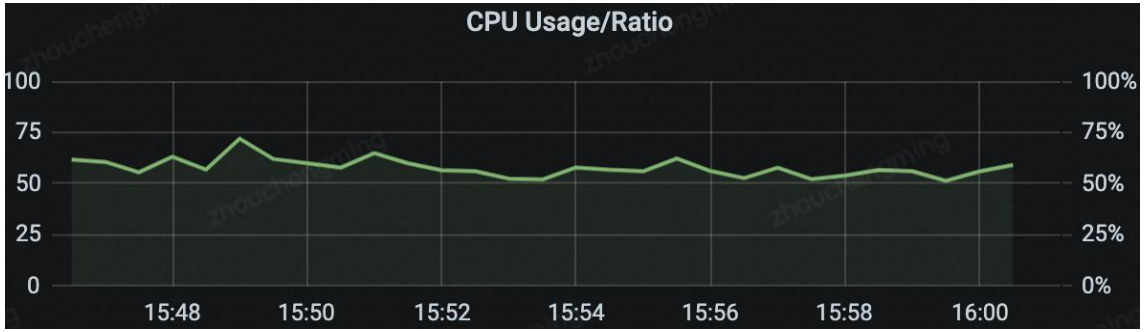

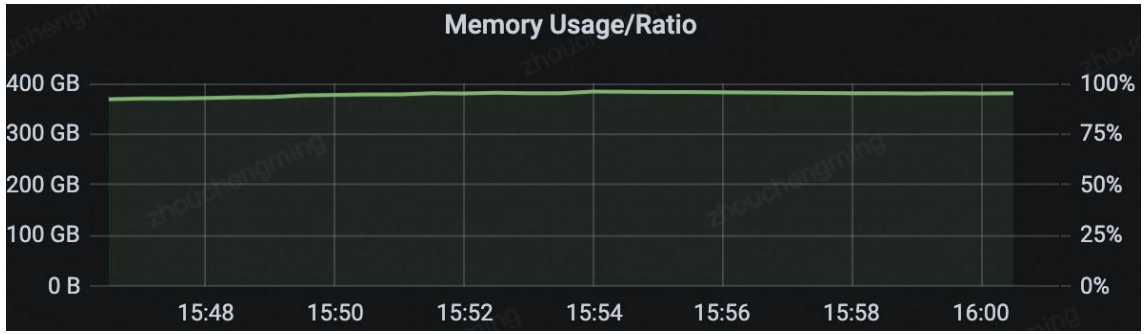

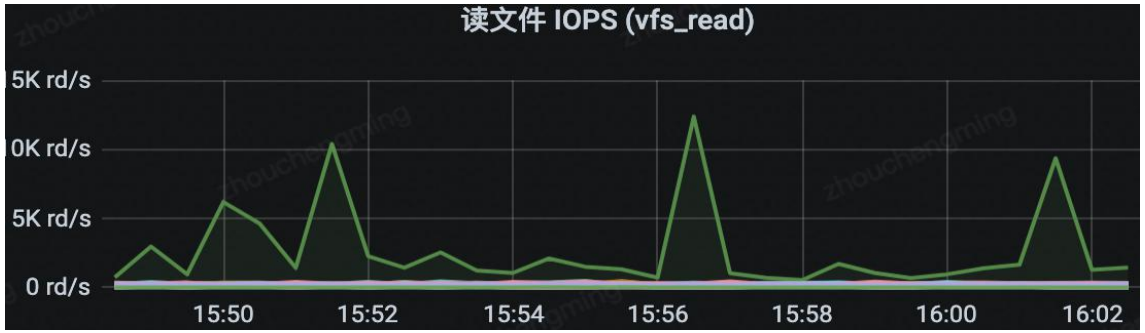

Cgroup

More Workload Control Groups

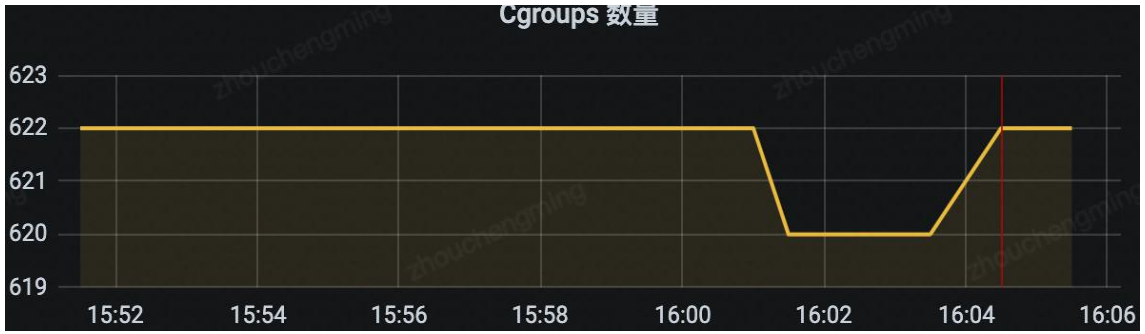

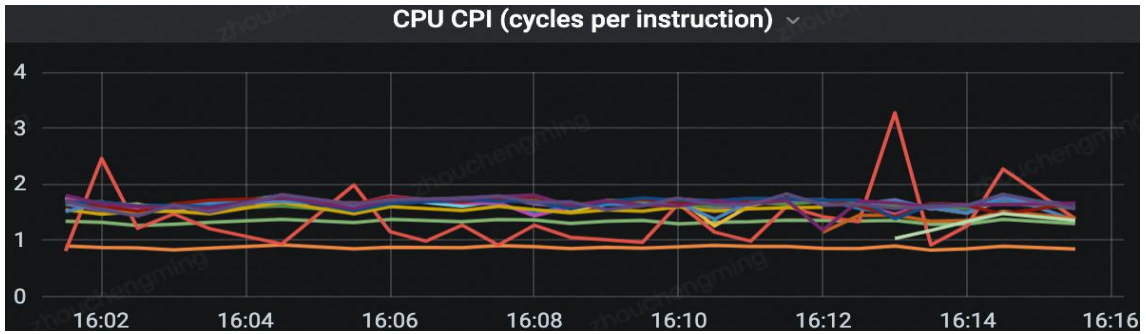

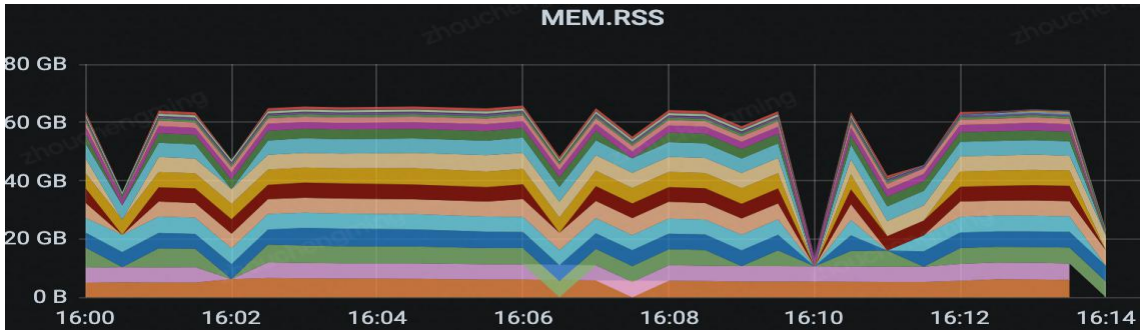

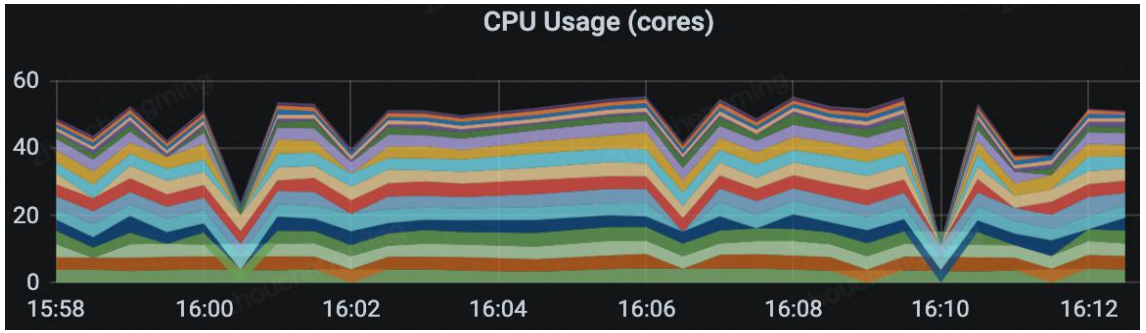

Pressure

- How much pressure are workload causing
  - Useful for performance debug like latency burst
  - Maximizing resource utilization while maintain performance
  - evict specified workloads to meet the needs of priority workloads
  - kill workloads to spare minimal resource for system usability

- Old ways
  - Load Average
  - Vmpressure

PSI introduction

ByteDance 字节跳动

Interface

- proc interface for system resource pressure
  - for CPU

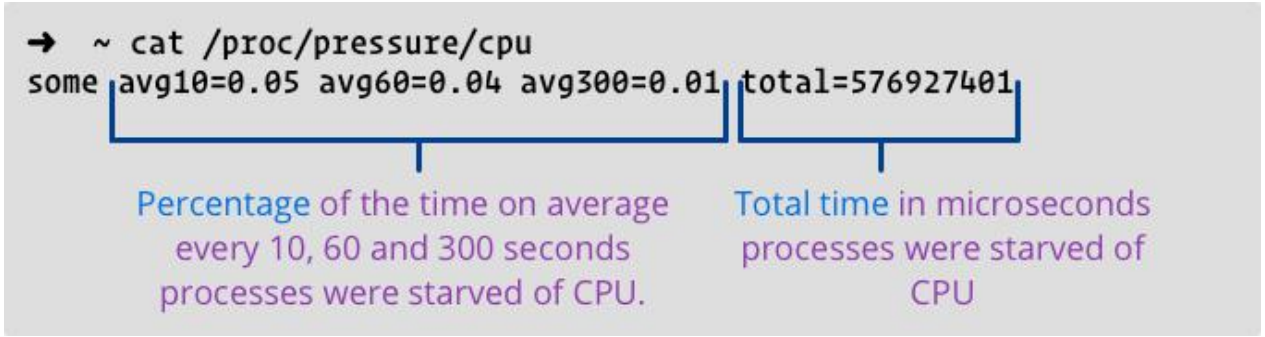

- cgroup interface for cgroup resource pressure
  - for memory and IO

    some avg10=0.00 avg60=0.00 avg300=0.00 total=0

    full avg10=0.00 avg60=0.00 avg300=0.00 total=0

Interface

- notification when reaching specified pressure threshold
  - echo “some 100000 1000000” > /proc/pressure/cpu
  - use poll/epoll/select to wait for notification

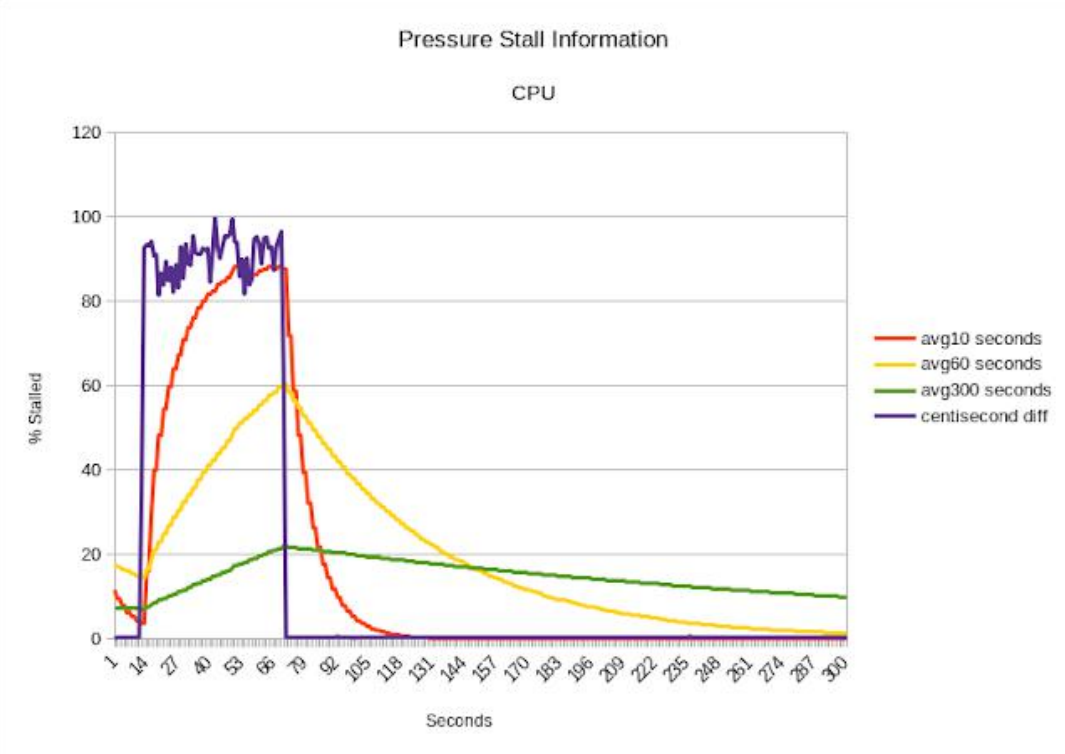

Definition

- some

  percentage of time some (one or more) tasks were delayed due to lack of resources

- full

  percentage of time in which all tasks were delayed by lack of resources

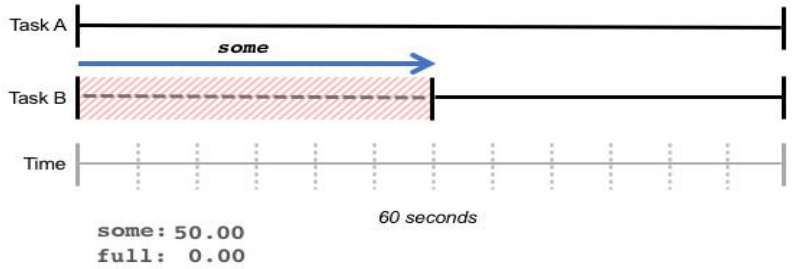

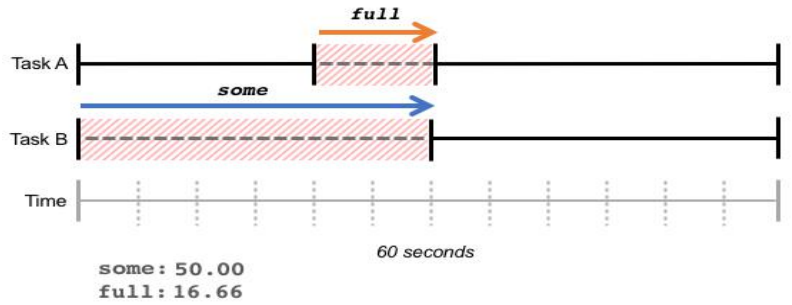

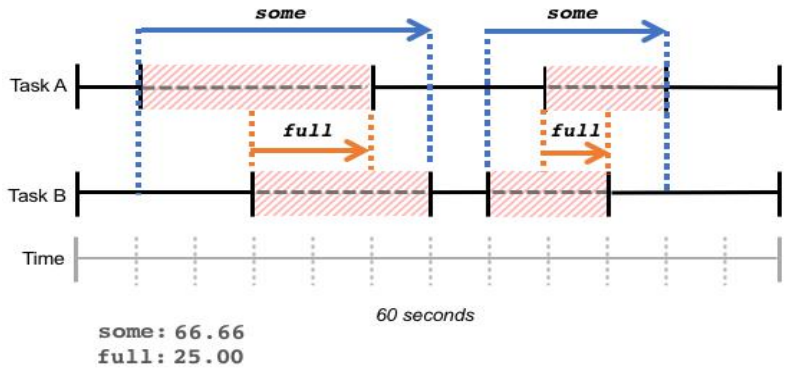

PSI improvement

ByteDance 字节跳动

Implementation

- Tracking task status
  - TSK_RUNNING
  - TSK_ONCPU
  - TSK_MEMSTALL
  - TSK_IOWAIT

- Accumulate task count per-cgroup per-CPU per-status
  - PSI_IO_SOME tasks[NR_IOWAIT] > 0
  - PSI_IO_FULL tasks[NR_IOWAIT] > 0 && task[NR_RUNNING] = 0
  - PSI_MEM_SOME tasks[NR_MEMSTALL] > 0
  - PSI_MEM_FULL tasks[NR_MEMSTALL] > 0 && tasks[NR_RUNNING] = 0
  - PSI_CPU_SOME tasks[NR_RUNNING] > tasks[NR_ONCPU]
  - PSI_NONIDLE tasks[NR_RUNNING] || tasks[NR_MEMSTALL] || tasks[NR_IOWAIT]

Improvement

- Accumulate task count per-cgroup per-CPU per-status
  - PSI_IO_SOME tasks[NR_IOWAIT] > 0
  - PSI_IO_FULL tasks[NR_IOWAIT] > 0 && task[NR_RUNNING] = 0
  - PSI_MEM_SOME tasks[NR_MEMSTALL] > 0
  - PSI_MEM_FULL tasks[NR_MEMSTALL] > 0 && tasks[NR_RUNNING] = 0
  - PSI_CPU_SOME tasks[NR_RUNNING] > tasks[NR_ONCPU]
  - PSI_CPU_FULL tasks[NR_RUNNING] > 0 && tasks[NR_ONCPU] = 0
  - PSI_NONIDLE tasks[NR_RUNNING] || tasks[NR_MEMSTALL] || tasks[NR_IOWAIT]

- Meaningful at the cgroup level
  - means all non-idle tasks in a cgroup are delayed on the CPU resource
  - throttled by the CPU bandwidth control
  - CPU used by other cgroups

Tracking

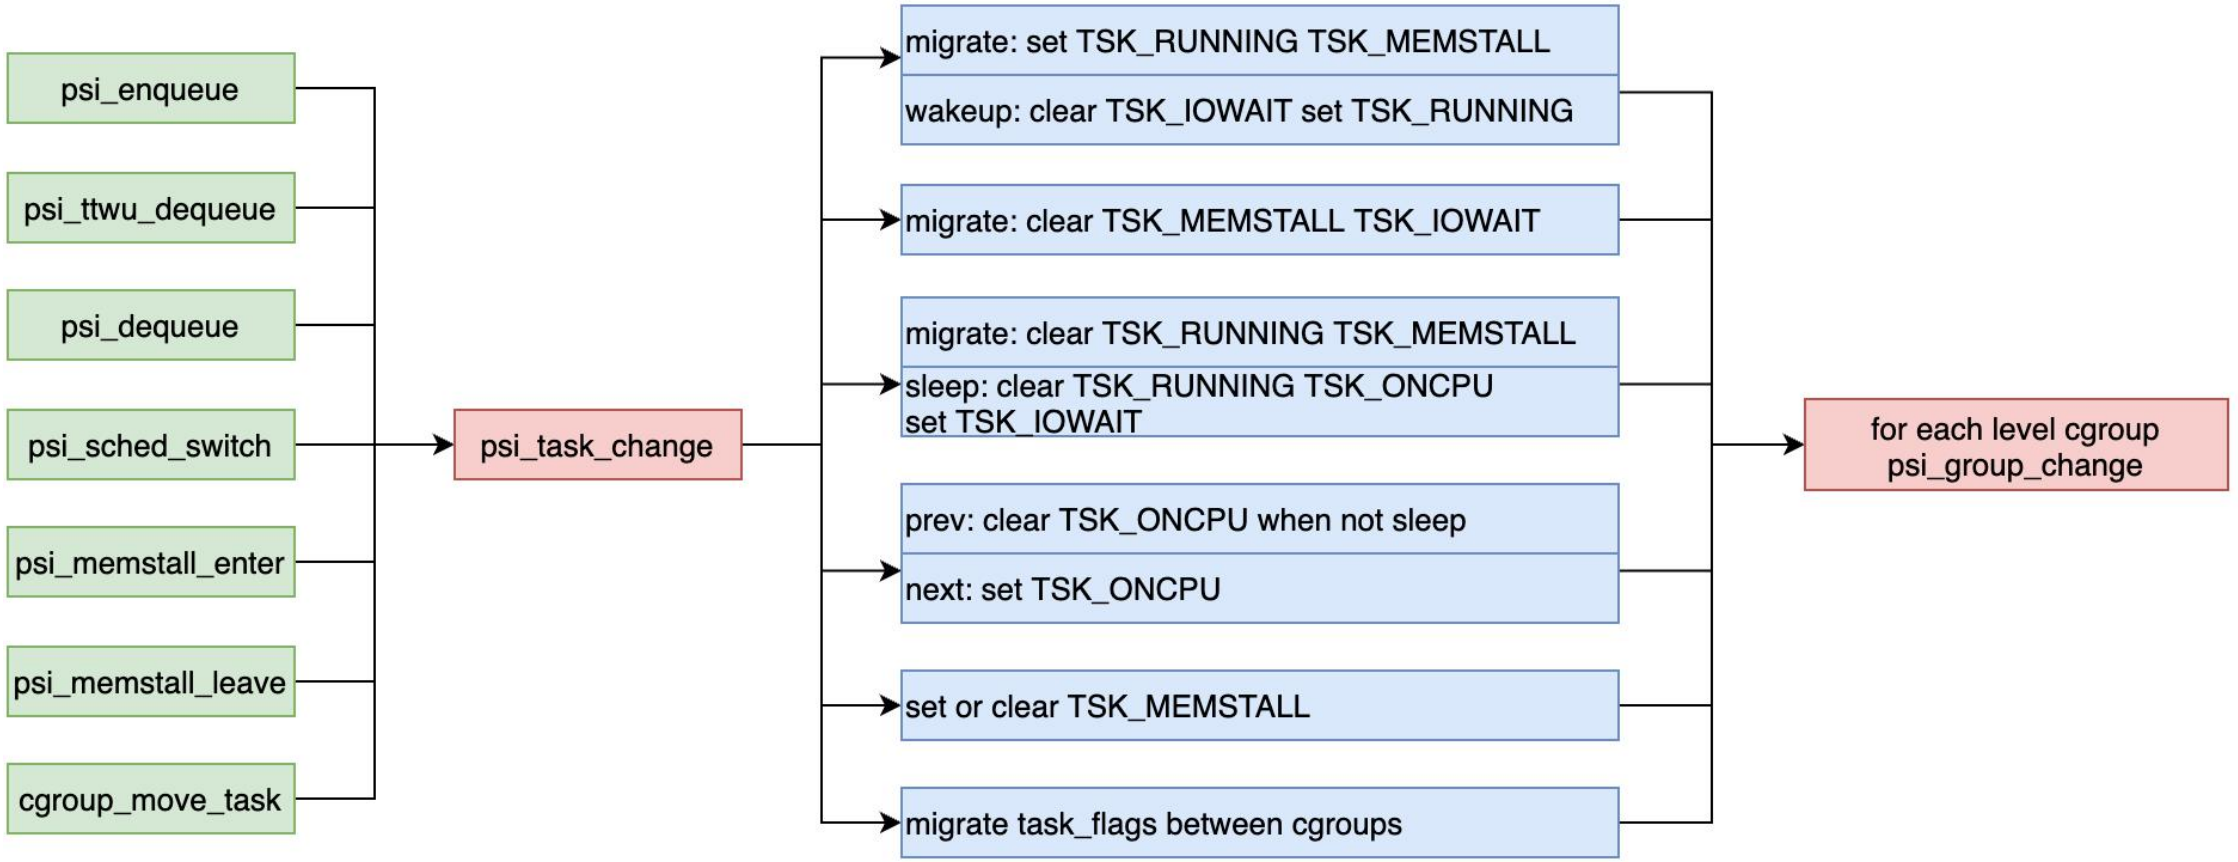

Improvement

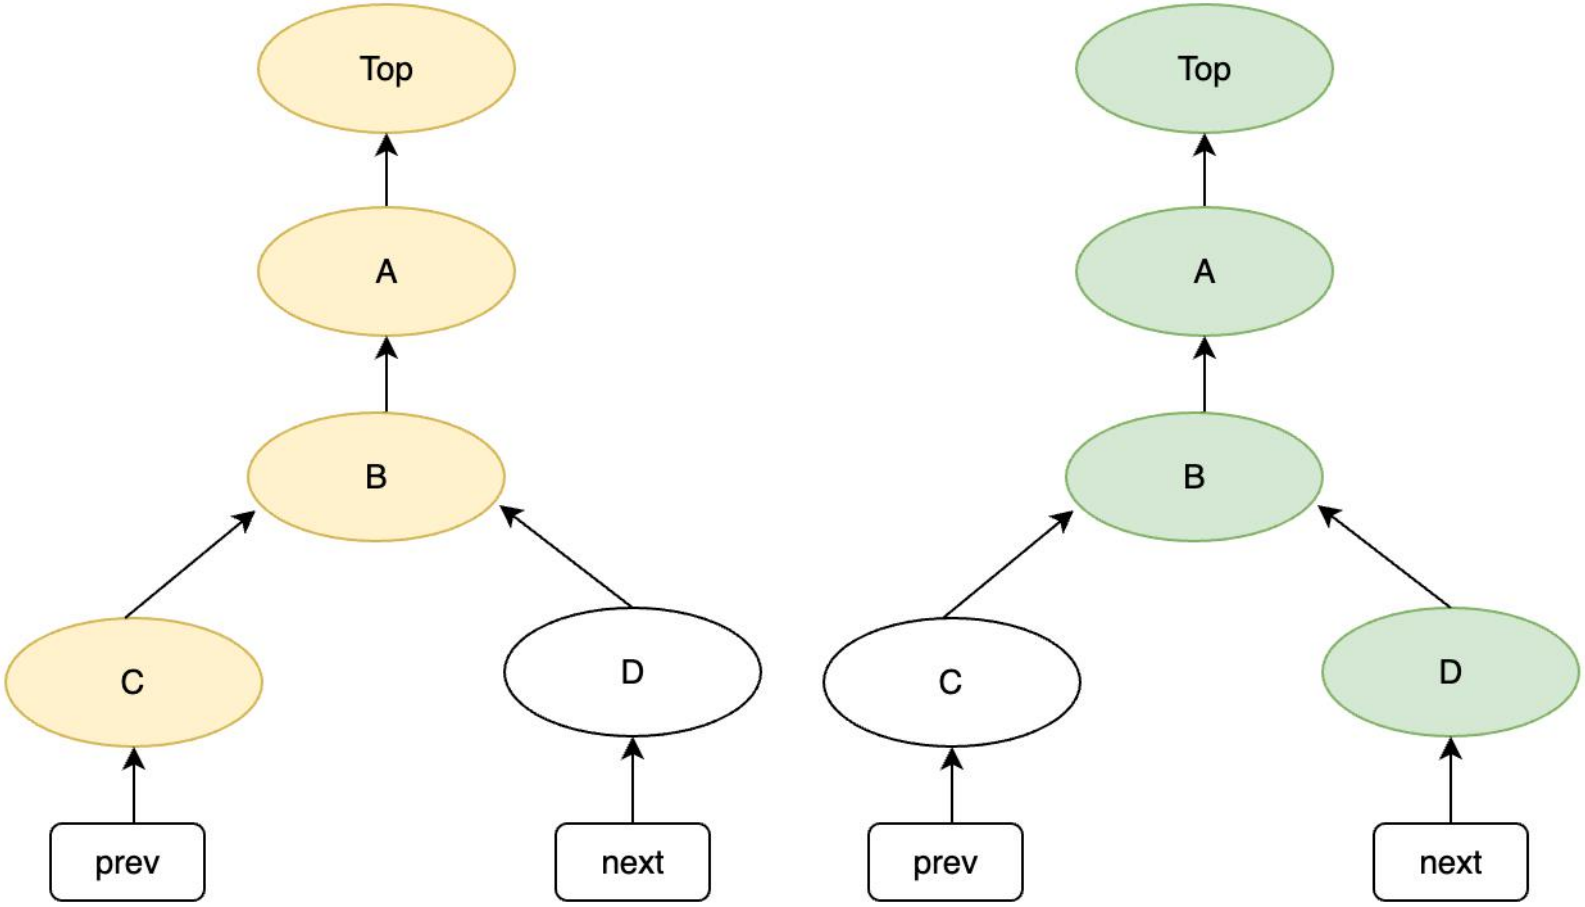

Improvement

- task switch in deep cgroup levels
  - clear prev TSK_ONCPU until Top
  - set next TSK_ONCPU until Top
- psi_task_switch preempt case
  - clear prev TSK_ONCPU until B
  - set next TSK_ONCPU until B
- psi_task_switch sleep case
  - delay psi_dequeue to psi_task_switch
  - avoid ONCPU changes on common ancestors

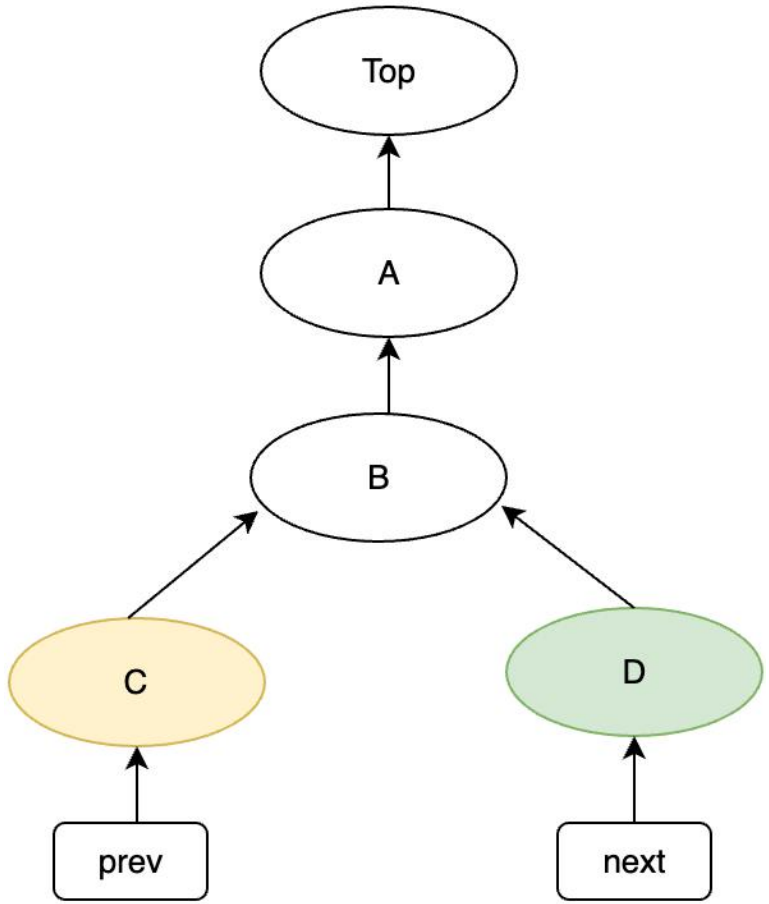

Improvement

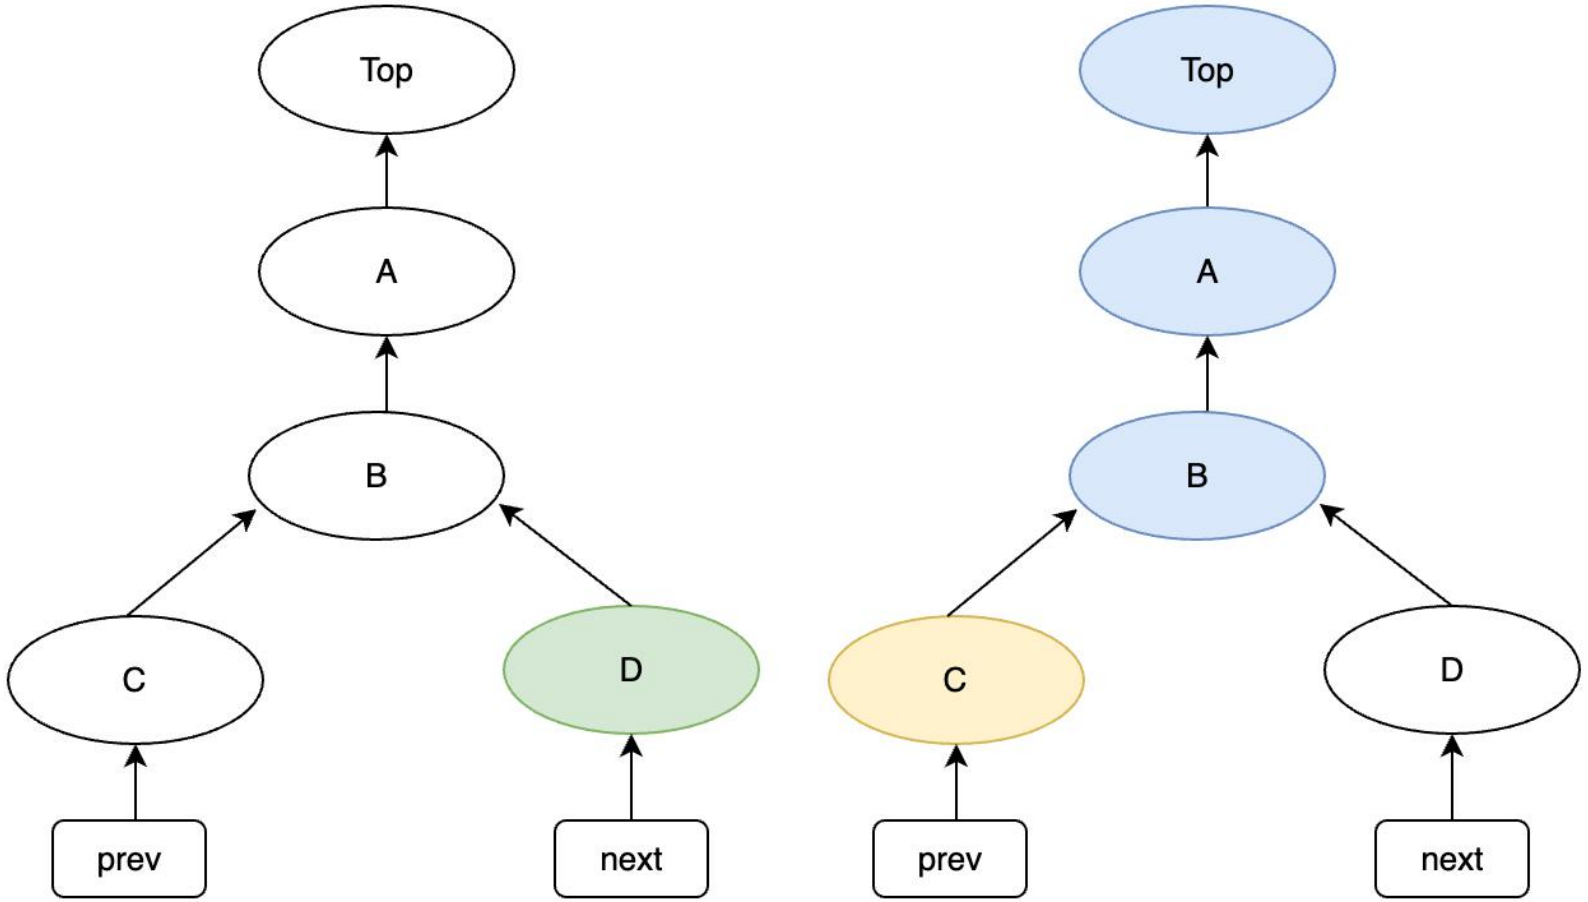

Improvement

- sleep before:
  - psi_dequeue()
    ```
    while ((group = iterate_groups(prev)))  # all ancestors
        psi_group_change(prev, .clear=TSK_RUNNING|TSK_ONCPU)
    ```
  - psi_task_switch()
    ```
    while ((group = iterate_groups(next)))  # all ancestors
        psi_group_change(next, .set=TSK_ONCPU)
    ```

- sleep after:
  - psi_dequeue()
    ```
    nop
    ```
  - psi_task_switch()
    ```
    while ((group = iterate_groups(next)))  # until (prev & next)
        psi_group_change(next, .set=TSK_ONCPU)
    while ((group = iterate_groups(prev)))  # all ancestors
        psi_group_change(prev, .clear=common?TSK_RUNNING:TSK_RUNNING|TSK_ONCPU)
    ```

Gain

The CPU overhead of PSI is reduced by about 10% for our workload.

Status & Future Work

ByteDance 字节跳动

Status & Future Work

Status

- Improvement patchset has been merged into Linux kernel.
  - https://lore.kernel.org/lkml/20210303034659.91735-1-zhouchengming@bytedance.com/

Future Work

- More task status and more PSI metrics
- Monitor and notification for cgroup adaptive tuning
- More performance improvements in task and cgroup tracking

THANKS.

ByteDance字节跳动