

在虚拟化场景中构建 基于硬件的性能监控服务

Like Xu

腾讯云虚拟化开源团队

CLK 2021



Agenda

1. A basic **mindset** of CPU performance analysis
2. **Perf** profiling modes and its subcommands roadmap
3. Current publicly **available** hardware capabilities on x86
 1. The **PMC** workflow and some devilish details
 2. Virtualizing PMC and the basic **KVM framework**
 3. Virtualizing **Branches** Sampling Facilities
 4. Virtualizing Instructions **Trace** Facilities
4. Challenges in the ongoing **hybrid** scenarios
5. (Virtualized) PMU **Use Cases** at Tencent Cloud

1. A basic mindset of CPU performance analysis

- Optimization is driven by careful performance analysis, not intuition.
 - Full stack, short board, joggle, spikes, long tail
 - Telemetry/APM agent → Please fix other issues first
- How can we further accelerate our code ?
 - HW utilization and saturation, and also errors.
 - even if the CPU occupancy (not real utilization) reaches 99%
- Increase as many Instructions per cycle (IPC) ↑ as possible → **Optimizable**

Architecture
HW events

- elapsed time ↓, retired instructions per functional block ↓
- real walked cycles (not crystal clock/tsc) before the next branch ↓
- cache-references ↑ + cache-misses ↓
 - cache (L, D, TLB) ↑ and memory latency ↓ hierarchy
 - branch-instructions ↓ + branch-misses ↓
 - frontend stalled cycles ↓
- utilization of execution units ↑
 - backend stalled cycles ↓
 - execution units and types (Int, FP, Scalar, Vector) ↑

Model Specified
HW events

- More CPU HW details
 - Out of order Window, Scheduler Entries ...
 - Register Files, Allocation Queue ...
 - Smarter speculation/prefetch ... algorithms
 - Cstate change, uncore iio/imc ...
 - check <https://perfmon-events.intel.com/>

FE → I-Cache, Decode
Branch, μop cache

BE → Mem load/store
Execution, D-Cache

Is μOP
Queued
And
Retired
?

Performance bottlenecks

Scalable Application Code

Distributed Services

Components and language

Runtime System (Lib, OS, VM, MM)

Data structures, algorithms

Instructions Branch Blocks

HW Platform Configurations
(dies, mem chips, firmware, power)

ISA (Architecture)

CPU Microarchitecture

Vendor Logic/Devices/Electrons

Observable

Monitoring interesting event with counter (PMC) →

How the instructions flow reaches this event →

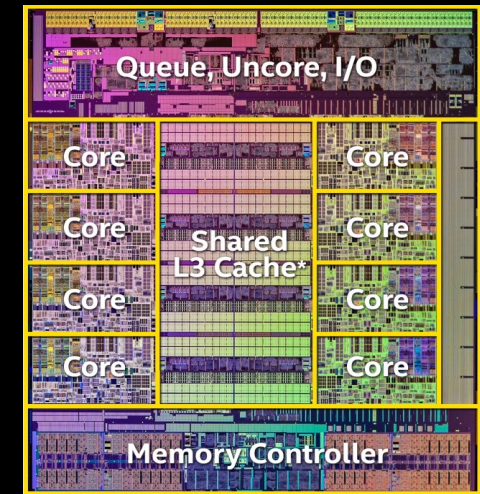
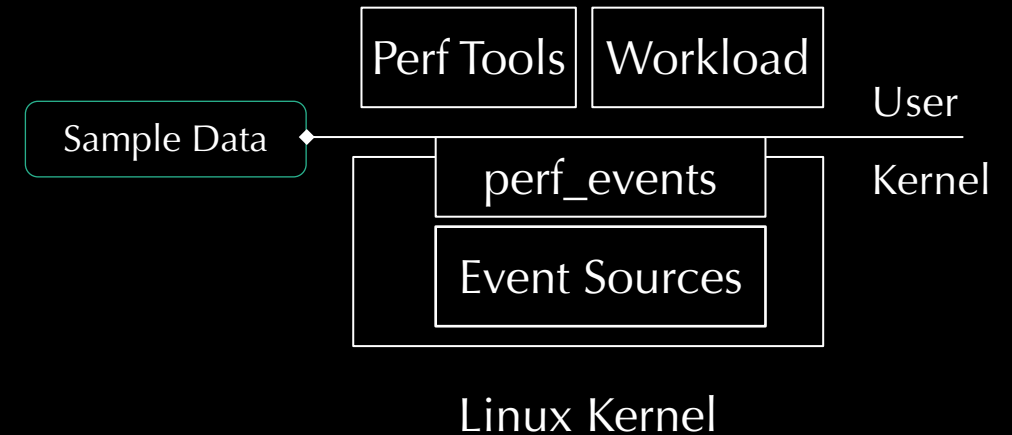
Monitoring hardware with hardware (PMU) →

Next Page : Perf - Linux Kernel official performance analyzing tool ←

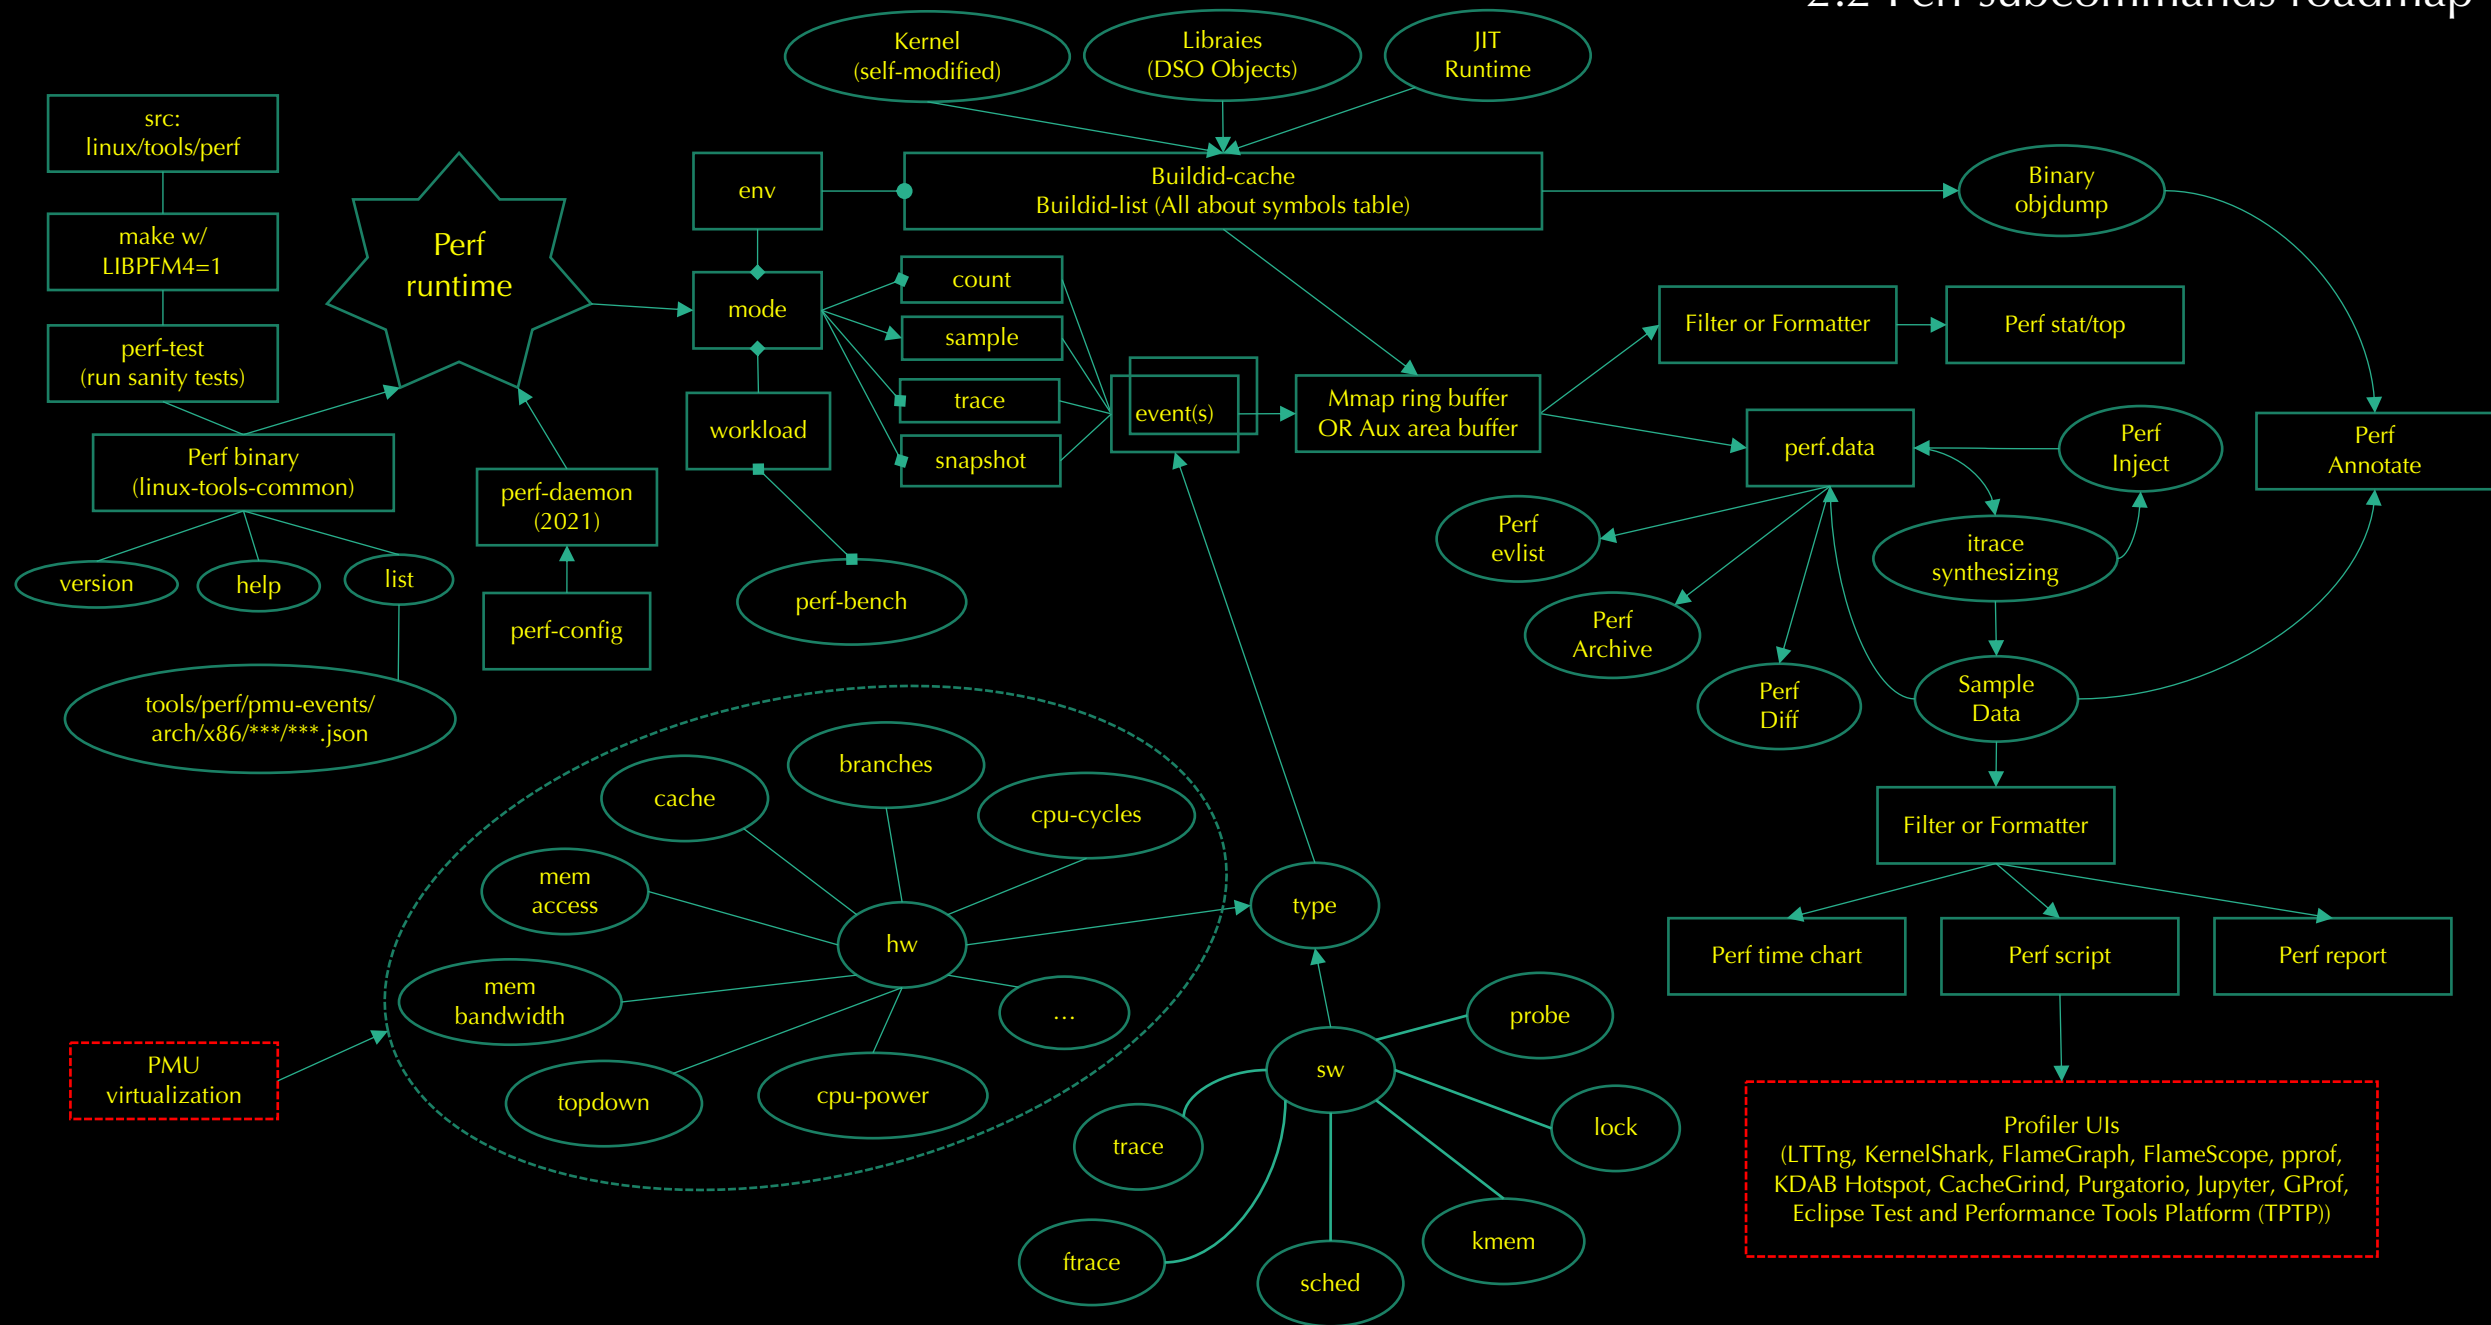
- Workload Classification :: Top-down Microarchitecture Analysis (not for CLK)
 - http://www.cs.technion.ac.il/~erangi/TMA_using_Linux_perf_Ahmad_Yasin.pdf
 - keep performance in mind as you make early design and architectural decisions

2.1 Perf profiling modes

- Profiling Target
 - system wide, per-user/numa/cpu/thread/
 - Namespaces/cgroup/dso(dynamic shared object (DSO))/
 - command/guest, exclude-perf
- Profiling Lifecycle :: Count -> Sampling -> Snapshot -> Trace -> Debug -> Count
- Count mode
 - `$ perf stat --interval-clear -I 1000 --metrics IPC -C 0`
 - `$ perf stat --topdown --td-level=2 --no-metric-only`
- Periodic Sampling mode --> `$ perf record ...`
 - based on the occurrence of a particular event such as a timer or interrupt
 - Record w/ -F, -c,
 - `$ perf top -e 'cache-misses,cycles' -v --force -F 49 --realtime=1 -u root \`
`-s overhead,comm --call-graph lbr --no-children --percentage relative --show-nr-samples`
- Full-trace mode
 - Get all execution records
 - Software trace
 - `$ perf trace/fttrace/probe/kmem/lock/sched/kvm ...`
 - Hardware trace (more on next slide)
 - Intel branch trace store/Processor Trace/LBR call-stack mode
- Snapshot mode :: let trace run and overwrite older data in the buffer
 - Run ring buffer, stop trace on event of interest, save only tail of trace
 - `$ perf record -v -e intel_pt//u -S ./loopy 1000000000 &`
 - `$ kill -USR2 11435 # Recording AUX area tracing snapshot`
- Mixed mode (such as, count + sample)
 - `$ perf stat record ...`



2.2 Perf subcommands roadmap



2.3 Sample Data Fields

```
$ perf script -F comm,tid,pid,cpu,time,event,ip,dso,sym,symoff,srcline,period,insnlen,insn,misc
```

```
CPU 115/KVM 466675/466796 [076] K      106883.295520:      35940   cycles: ffffffff81e31d40 vmx_vmexit+0x0 (/lib/modules/5.15.0-rc2/build/vmlinux)
vmenter.S:79 ilen: 2 insn: eb 30
```

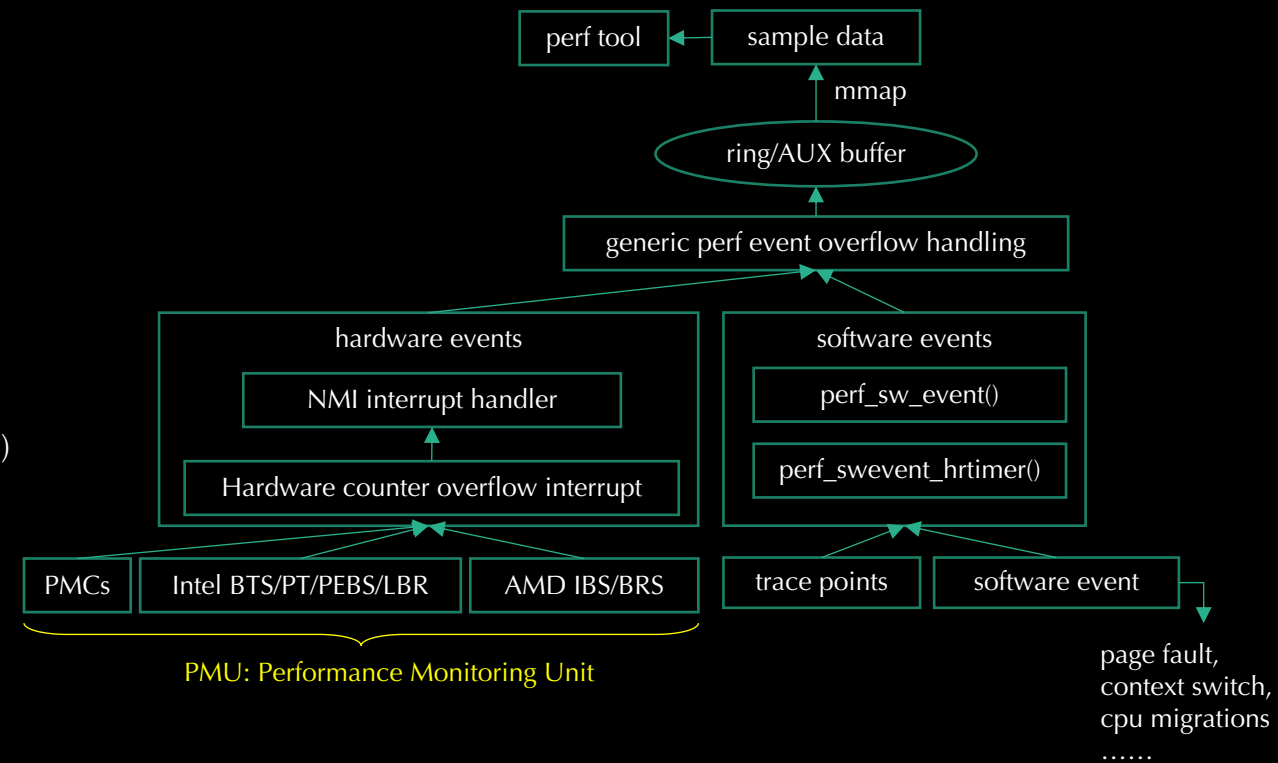
	comm,tid,pid,cpu,time,event,ip,dso,sym,symoff,srcline,period,insnlen,insn,misc	flags ,ipc,callindent # -e intel_pt//, --itrace=cr -F +callindent
branches: tr strt		addr # sample address, --data
branches: call	IPC: 0.07 (124/1637)	
branches: tr strt	IPC: 0.03 (17/475)	
branches: tr strt	ABI:2 AX:0xff BX:0xff11003f772137e0	tod # --clockid CLOCK_MONOTONIC
lock_acquire	ABI:2 AX:0xff BX:0xff11003f772137e0	iregs, # --intr-regs
rcu_read_lock_sched_held	ABI:2 AX:0xff BX:0xff11003f772137e0	uregs, # --user-regs
__fentry__	2M	data_page_size, # -e "mem-loads,mem-stores" --data-page-size
__mem_cgroup_uncharge_list	1024M	code_page_size # -e "mem-loads,mem-stores" --code-page-size
uncharge_page	4K	phys_addr # -e "mem-loads,mem-stores" --phys-data
__fentry__		brstack # -b
release_pages	fffffffff811f8267 10e5d0e9a	brstacksym,brstackoff # -b
free_unref_page_list	fffffffff811d3003 3f77fedc08	brstackinsn # -b
__fentry__	fffffffff811d3031 3f7f7ed18c	
tlb_flush_mmu		
tlb_flush_mmu	perf_event_exec+308:	
tlb_finish_mmu	fffffffff81359254 insn: 48 8b 7c 24 20	
__vma_adjust	fffffffff81359259 insn: e8 e2 fc ae 00	# PRED 9 cycles [260] 0.22 IPC

2.4 More lesser known tips on perf usages (a.k.a, perf tips Brendan Gregg won't tell you)

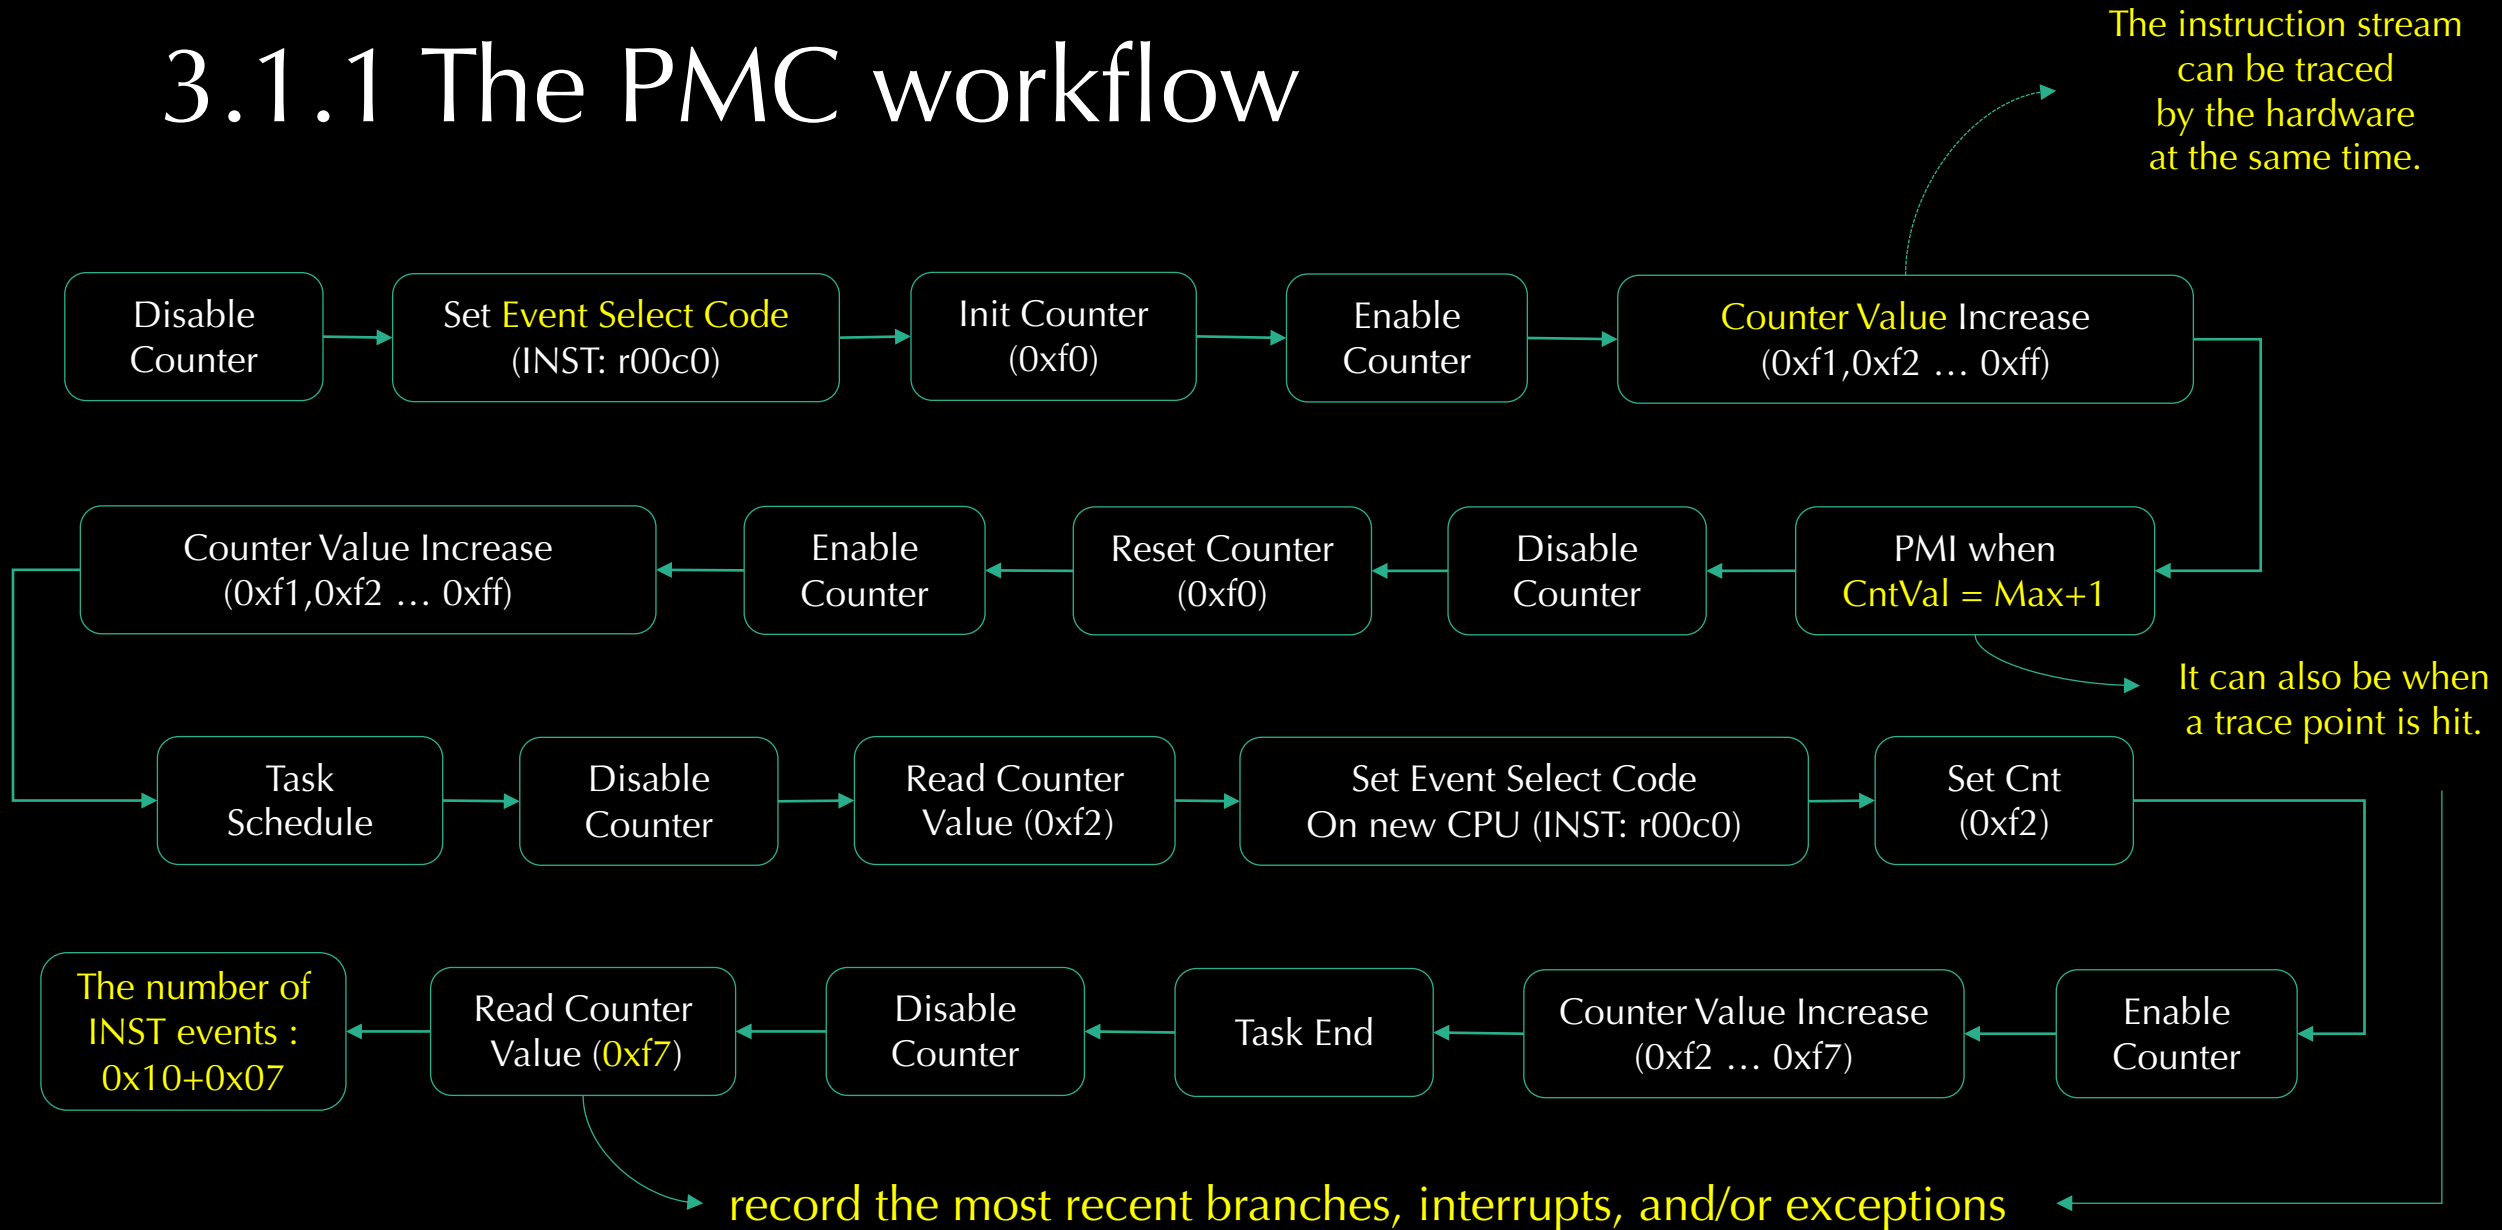
- How to build Perf in the best way
- How to run sanity tests for perf
- How to list all events
 - `$ perf list --desc --long-desc --details (-raw-dump)`
- Event switch-on/off tracing mode
- How to use uncore hardware events
 - `$ perf iostat '0000:16,0000:20' -l 1000`
- How to detect cache line false sharing w/ perf-c2c
- What is perf buildid & buildid-cache
- Many ways to program events
- Profile using event groups sampling
- What is advanced metric
 - `$ perf list metric`
- How to save copies of /proc/kcore, /proc/kallsyms and /proc/modules
- How to control mmap buffer behavior and the size of output perf.data
 - `--no-buffering, --mmap-pages, --mmap-flush, --overwrite`
- How to dump event raw sample/trace
- How to check each trace/sample
- How to synthesize samples from traces
- How to understand perf-annotate
- How to run perf session on background
- Tips to reduce perf record overhead
 - Multi AIO trace writing
 - Set affinity mask of trace reading thread
 - Produce compressed trace
 - `--proc-map-timeout`
- Perf built-in benchmark for kernel
- How to use perf-trace/perf-ftrace/perf-kmem/perf-lock/perf-kmem/perf-sched/perf-config ...
- No time to cover all this time but the slide is already available at https://github.com/xuliker/kde/raw/master/likexu_slides/Deep%20dive%20into%20Linux%20perf%20tool.pdf

3. Current publicly available hardware capabilities on x86

- Why do we need hardware assistance ?
 - Achieving **low cost and high resolution** profiling is an art.
 - Do not require recompilation, restart, or embedded instrumentation (though support it if present).
 - Run fast, non-intrusively, to minimize overhead and gross disruption
 - Measure really granular things
- Performance Monitoring Counters (Available)
 - Architecture events
 - Specific Topdown/PERF_METRICS counter on Intel
- Intel features (Upstream ed-or-ing)
 - Last **Branch** Recording ((Arch, Timed) LBR)
 - Processor Event Based Sampling ((Extended, Adaptative) **PEBS**)
 - Branch Trace Store (**BTS**)
 - Processor Trace (**PT**)
- AMD features (Under Development)
 - Instruction-Based Sampling (**IBS**)
 - Branch Sampling Feature (**BRS**)
 - Lightweight Profiling (LWP, for user space)
- More and more PMU hw features will be available.
 - Some inspiration for RISC-V design

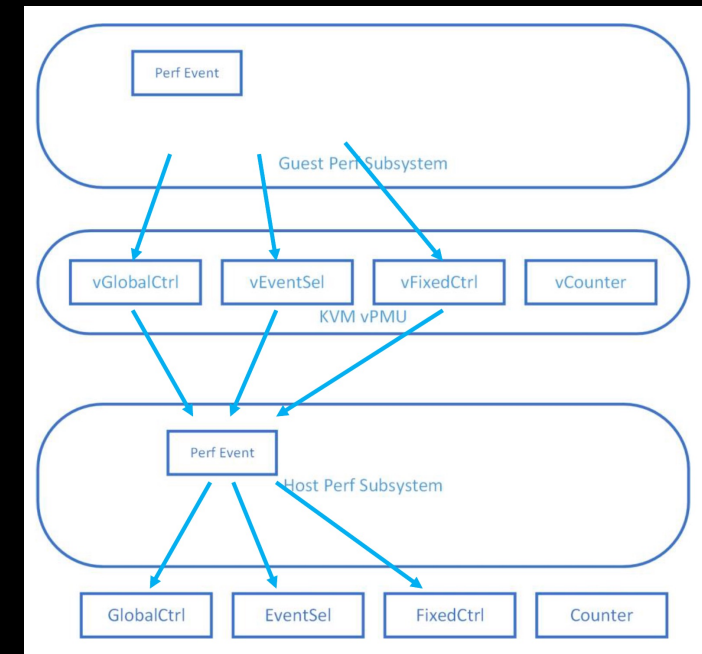


3.1.1 The PMC workflow



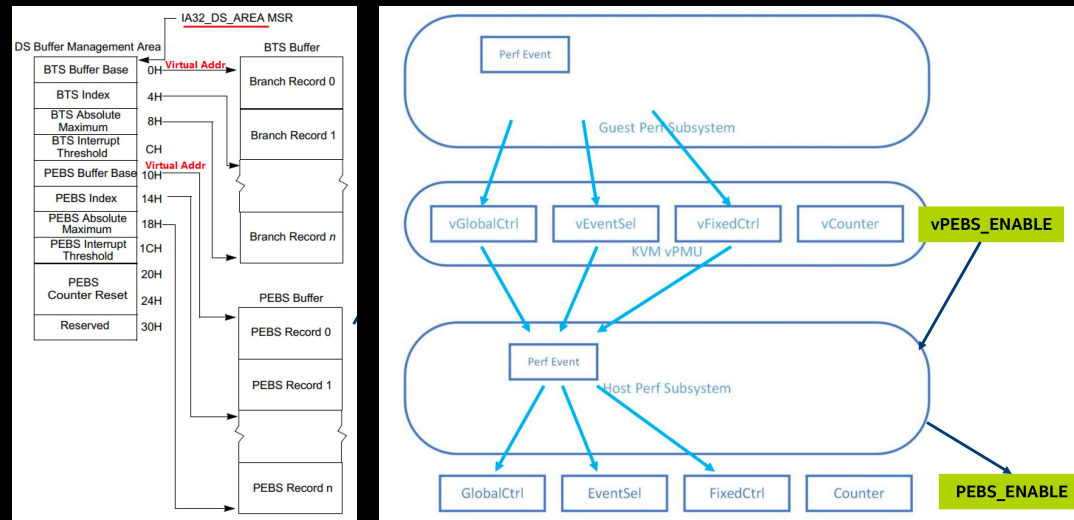
3.2 Virtualizing PMC and the basic KVM framework

- `$ perf record -e cpu/event=0xa8,umask=0x1,cmask=0x1/ ...`
- Emulating guest MSR-based counter resource
 - Event Code = Umask + EvtSel (per logical core)
 - RD/WR Event counter value
- Emulating guest counter enable
 - and disable behavior
- Emulating PMI for guest
- Emulating RDPMC for guest
 - Read Event Counter Value
- How to count for guess instruction events ?
- The scheduling of counters for guests emulation and host itself.



3.2 Virtualizing PEBS and the upstream story

- `$ perf mem --data-page-size record -e instructions:P ...`
- After the counter overflows,
 - the processor copies the current state of the general-purpose and EFLAGS registers and instruction pointer into a record in the precise event records buffer.
- The debug store (DS) buffer
- When the precise event records buffer is nearly full, an interrupt is generated



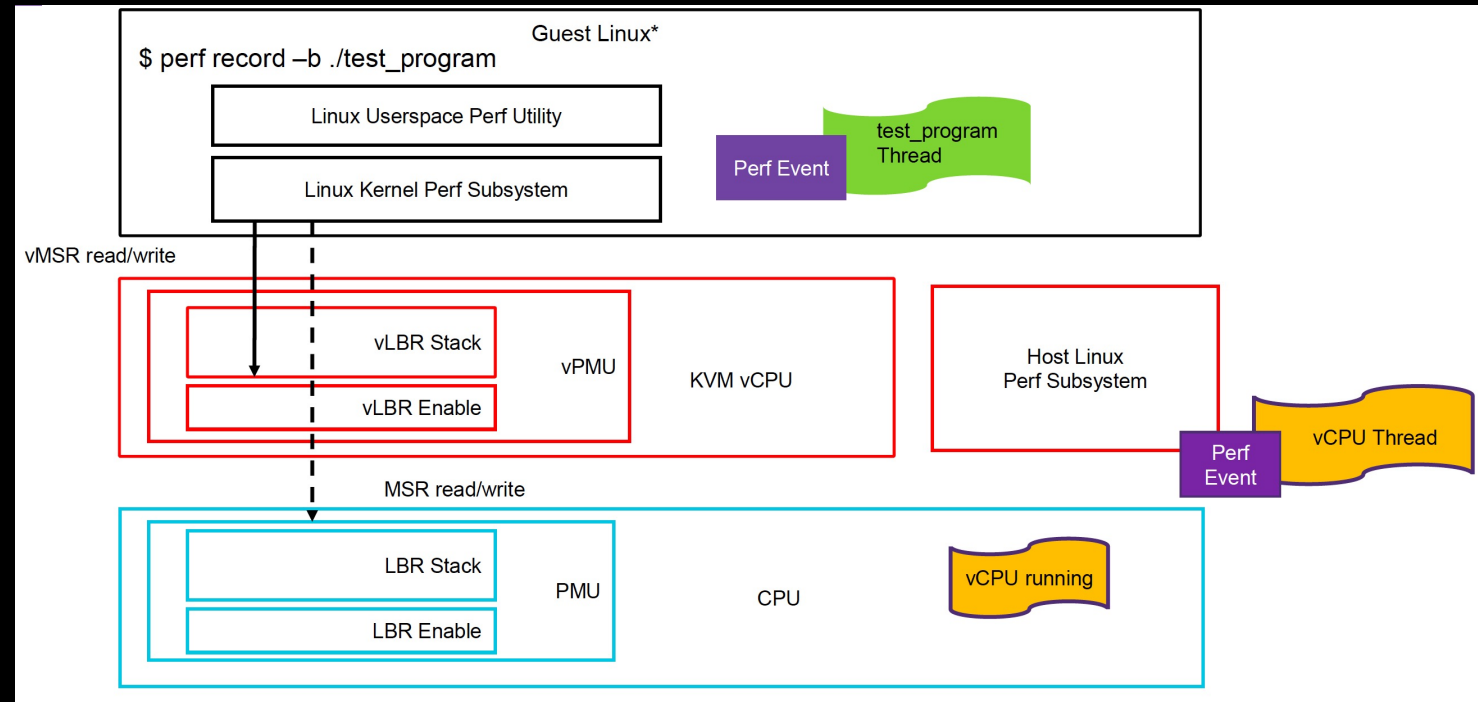
3.3 Virtualizing Branches Sampling Facilities

- **\$ perf record -g ... OR perf record --call-graph lbr ...**
- logging taken branches and control flow transfers within registers.
 - Intel LBR vs AMD BRS (Zen 3,19h)
 - Diff on lbr_depth, enable after counter overflow, support call_stack mode, hardware branch type filtering, fixed sampling period
 - enable autoFDO-style optimization by compilers
 - doesn't support LER (Last Event Record) due to security leakage

Branch Type	Operations Recorded
COND	Jcc, J*CXZ, and LOOP*
NEAR_IND_JMP	JMP r/m*
NEAR_REL_JMP	JMP rel*
NEAR_IND_CALL	CALL r/m*
NEAR_REL_CALL	CALL rel* (excluding CALLs to the next sequential IP)
NEAR_RET	RET (0C3H)
OTHER_BRANCH	JMP/CALL ptr*, JMP/CALL m*, RET (0C8H), SYS*, interrupts, exceptions, IRET, INT3, INTn, INTO, TSX Abort, EENTER, ERESUME, EEXIT, AEX, INIT, SIPI, RSM, breakpoints

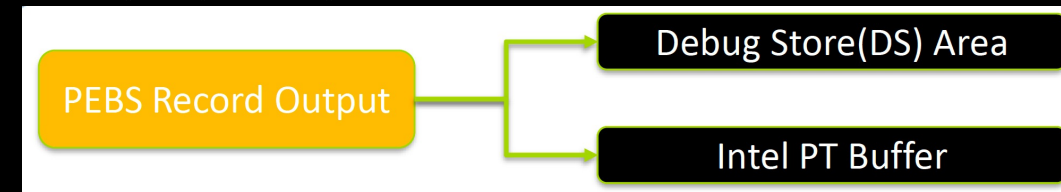
3.3 Virtualizing Branches Sampling Facilities

- Enabled via IA32_LBR_CTL or DEBUGCTRL MSR
 - VMX entry-load/exit-clear control
- What if host/guest branch record filters options are different
- LBR MSRs pass-through
- XSAVES/RESTORES
 - Speed up LBR read/reset
 - For Arch LBR



3.4 Virtualizing Instructions Trace Facilities

- Branch trace store (BTS) → just trace all branches taken
 - `$ perf record --per-thread -e intel_bts// ...`
- Intel PT :: log information about software execution
 - `$ perf record -e intel_pt// ...`
 - Supports control flow tracing
 - determine exact flow of software execution from trace log
 - with minimal impact to system execution (according to Intel)
 - Target <5% performance overhead
 - store both cycle count and timestamp information
 - Gdb 7.10 supports PT for “backwards debugging” (reverse-step)
 - ARM Coresight
- How to virtualizing PT
 - Intel PT VMX improvements will treat PT output addresses as **Guest Physical** Addresses (GPAs) and translate them using EPT that serves to simplify the process of Intel PT virtualization for using by a guest software.



4. Challenges in the ongoing hybrid scenarios

- Performance is a stabilizing feature and of its own right.
 - Performance on the Horizon
 - Zoom In/out
 - We need to satisfy any perf user in any profiling environment.
- Guest user need self-profiling its workload without any risk
 - How to keep both guest and host accurate
 - How to reduce guest and host profiling overhead
- For bare metal, system-wide profiling
 - Kernel space, user space, SMM space, even firmware
 - The tracking framework of the user space is quite colorful
- For CSP
 - Profiling virtual machine(s) and host at the same time
 - Collaboration and integration of guest kernels and hypervisor
 - `$ perf kvm --host --guest stat/record ...`
- For cloud native
 - service layer framework & business customized code
- In the near future
 - TEE (trusted execution environment) code & Non-TEE code
 - Dance with heterogeneous profiling systems facilities
 - eBPF program may need to touch PMU hardware
 - Heisenberg, aka observer effect

5. (Virtualized) PMU Use Cases at Tencent Cloud

- Dedicated Infrastructure Performance Optimization (not SRE) teams
- Performance debugging and code optimization for various languages
 - IPC ↑ forever, hot/critical path, frame pointerless stacktrace/callgraphs
 - Continuous compiler feedback, machine code layout optimizations
 - Garbage collection latency (JVMs, with GC cycles), lockless algorithms, zero-copy optimizer
 - TSX lock elision, deadlock, cache line false sharing
- Performance characterization for cloud workloads
 - NUMA/NVDIMM cold/hot page migration policy
 - Detect whether a virtual machine is live migration friendly
- Running perf as an part of program (libperf, perfmon2-libpfm4, eBPF)
- One of telemetry sources, runtime remotely triggering/configuration/instrumentation/collection/visualization
- Subscribe mode for uncore events
 - ask profiling data from various agents
- Performance Isolation Verification (such as RDT/cgroup v2)
- Observing the CPU health of both host and services
 - Add performance regression report to the DevOps pipeline, sub-second level of resolution
 - Security Intrusion (DDoS) Detection
 - Avoid causing the application to violate deadlines/SLA/QoS
- Analyze past, debug present, monitor future
 - Add performance regression to cover all vendor code changes
 - Kernel, Compilers, Retpoline/trampoline, Hotfixes, QEMU, Unikernel, virtio-fs, OCI runtime, block layer



Q & A

Thank you to join me at CLK 2021!
Error correction will be placed [here](#).

Like Xu

